Chenyu Li

chenyu157@utexas.edu | +1 512 647 0896 | Homepage

Education

The University of Texas at Austin Aug. 2026–Aug. 2028 (Expected)
B.S. in Informatics, Human-Centered Data Science Concentration

University of Liverpool & Xi'an Jiaotong-Liverpool University Sept. 2024–Jun. 2026 (Transferred Out)
BSc Information and Computing Science *GPA: 3.92/4.00*

- **Relevant Coursework**: Linear Algebra, Calculus, Discrete Mathematics, Data Structures, Algorithm and Problem Solving, Artificial Intelligence, Computer Systems, Operating Systems.

Research Interests

My research interests focus on the capabilities, safety, and societal implications of large language models. Specifically, I am interested in **long-horizon AI agents** for complex tasks, including **scientific discovery** and **AI-for-AI research**. I also study **responsible AI and computational social science**, focusing on AI safety, alignment, and the impact of AI on people and society. Furthermore, I explore **interdisciplinary AI research**, especially the integration of AI with fields such as law and psychology. Finally, I focus on **LLM evaluation** across these directions to assess model capabilities, safety, alignment, and performance in diverse contexts.

Publications

ScopeBench: Evaluating Cultural Norm Scope Awareness in LLMs
Haojun Liang*, Ziwei Zhao*, **Chenyu Li**, Qiyue Zhao, Rongduomijia Mu, Yulan Shen, Yu Gu, Chen Zhang, Kairui Zhang, Yanbiao Ma, and Wenxuan Wang. *EMNLP 2026 Main Conference*, pp. 27929–27960. (* Equal contribution.)

Manuscripts in Submission

Employing Large Language Models to Simulate Human Behaviors in Agent-Based Social Science Simulations
Xinnong Zhang, Xinyi Mou, Jia Wang, **Chenyu Li**, Shijun Lei, Jianing Shi, Zijian Ling, Zhenfei Yin, and Zhongyu Wei. *Manuscript in submission.*

Research Experience

Research Assistant | Tsinghua University, THUNLP Jul. 2026–Present
Advisors: Yuxiao Ye, Huiyuan Xie, and Zhiyuan Liu

- **CoGaBench** A theory-grounded benchmark evaluating LLMs' game-theoretic competence and Theory of Mind through strategic games. Primarily designed controlled evaluation games and probes targeting belief modeling, strategic planning, and adaptive decision-making.
- **Lexscope-CN** Lead a three-level benchmark of LLMs' understanding of Chinese legal norms: rule knowledge, recognition of nationwide versus province-specific scope, and application to concrete cross-regional legal cases.

Research Assistant | University of Oxford, Department of Engineering Science Oct. 2025–Sept. 2026
Advisor: Zhenfei Yin

- **SocialScientistBench** Co-developed 20+ Research Task Packages across 7 social-science disciplines in collaboration with domain experts, including paired End-to-End Tasks that contrast supported scientific competence with autonomous research. Built standardized-model and native-agent evaluation tracks and a task-specific plus eight-dimension scientific-quality framework to identify where autonomy breaks down.
- **LLM Agents in Agent-Based Modeling** Extended a Mesa-based architecture to support hybrid ABM and LLM-agent simulations. Designed Prisoner's Dilemma evaluations across 15 fixed strategies to test whether LLM agents follow assigned strategic behaviors, and analyzed their decisions in repeated mixed-agent interactions. Contributed to a manuscript on LLM-agent behavior in social-science simulations.

Research Assistant | Renmin University of China, AIM3 Lab Aug. 2025–May 2026
Advisor: Wenxuan Wang

- **ScopeBench** Led idea development and experimental design. Curated 3,000+ cultural rules across 50 countries and developed corresponding scope-judgment questions. Designed 5,000+ cultural-conflict scenarios to assess the cultural appropriateness of LLM-generated advice. Contributed to "ScopeBench: Evaluating Cultural Norm Scope Awareness in LLMs," published at EMNLP 2026 Main Conference.

Industry Experience

Research Intern | Microsoft Research Asia, Social Computing Group Mar. 2026–Jul. 2026
Advisors: Fangzhao Wu and Xing Xie

- Evaluated robustness to jailbreak and indirect prompt-injection attacks in collaboration with Rakuten AI, and tested prompt-based and SFT-based mitigation strategies.
- Developed a baseline pipeline with the Microsoft News product team for daily news push generation and ran real-world click-through experiments to measure user engagement.
- Contributed to a study of model safety degradation under incorrect-knowledge fine-tuning.

Honors & Awards

Stars of Tomorrow, Microsoft Research Asia 2026
Xi'an Jiaotong-Liverpool University Continuing Academic Merit Scholarship (¥93,000) 2025
Xi'an Jiaotong-Liverpool University Summer Research Undergraduate Fellowship (¥1,000) 2025
Xi'an Jiaotong-Liverpool University President's Entrance Scholarship (¥93,000) 2024

Skills

Programming Languages: Python, Java, C, MySQL
Development Tools: PyTorch, vLLM, LLaMA-Factory, Linux, Git, LaTeX
Languages: Chinese Mandarin (Native), English (Fluent)